

DATASET - IRIS

1. Data Analysis

Load the dataset and perform initial exploratory data analysis (EDA) to understand its structure and characteristics.

```
In [5]: # Import necessary Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [6]: # Load the Iris dataset
iris_df = pd.read_csv('IRIS.csv')
iris_df
```

```
Out[6]:
```

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa
...
145	6.7	3.0	5.2	2.3	Iris-virginica
146	6.3	2.5	5.0	1.9	Iris-virginica
147	6.5	3.0	5.2	2.0	Iris-virginica
148	6.2	3.4	5.4	2.3	Iris-virginica
149	5.9	3.0	5.1	1.8	Iris-virginica

150 rows × 5 columns

```
In [7]: # Display basic information about the dataset
print(iris_df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   sepal_length    150 non-null   float64
1   sepal_width     150 non-null   float64
2   petal_length    150 non-null   float64
3   petal_width     150 non-null   float64
4   species         150 non-null   object
dtypes: float64(4), object(1)
memory usage: 6.0+ KB
None
```

```
In [8]: # Summary statistics
print(iris_df.describe())
```

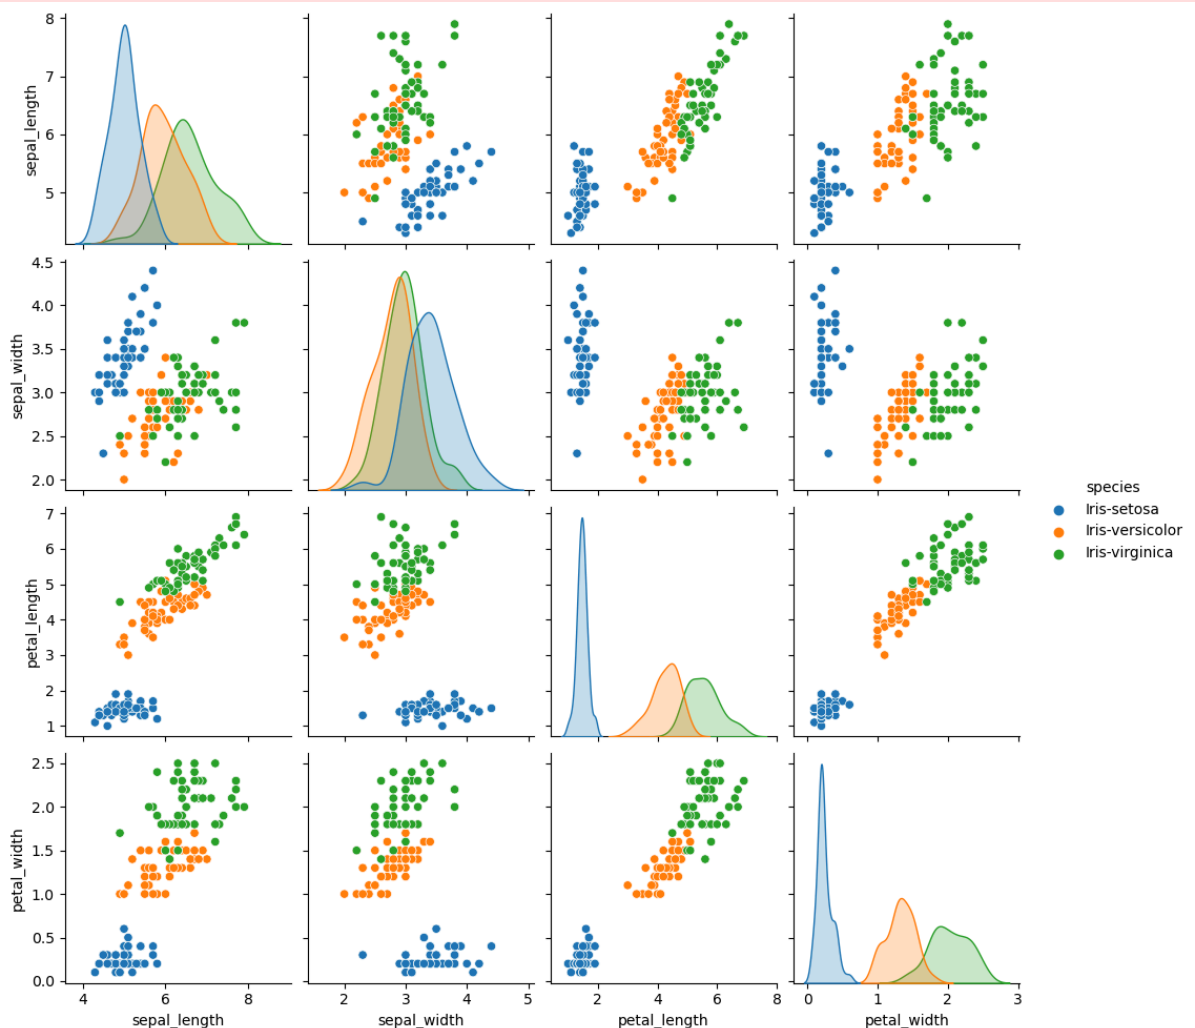
	sepal_length	sepal_width	petal_length	petal_width
count	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.054000	3.758667	1.198667
std	0.828066	0.433594	1.764420	0.763161
min	4.300000	2.000000	1.000000	0.100000
25%	5.100000	2.800000	1.600000	0.300000
50%	5.800000	3.000000	4.350000	1.300000
75%	6.400000	3.300000	5.100000	1.800000
max	7.900000	4.400000	6.900000	2.500000

```
In [9]: # Class distribution
print(iris_df['species'].value_counts())
```

```
species
Iris-setosa      50
Iris-versicolor  50
Iris-virginica   50
Name: count, dtype: int64
```

```
In [10]: # Pairplot for initial visualization
sns.pairplot(iris_df, hue='species')
plt.show()
```

C:\Users\Nishita Bala\anaconda3\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The figure layout has changed to tight
self.figure.tight_layout(*args, **kwargs)



2. Preprocessing

Prepare the data for modeling by handling missing values, encoding categorical variables, and splitting into training and testing sets.

```
In [11]: from sklearn.preprocessing import LabelEncoder  
from sklearn.model_selection import train_test_split
```

```
In [12]: # Encode categorical target variable 'species'  
le = LabelEncoder()  
iris_df['species_encoded'] = le.fit_transform(iris_df['species'])
```

```
In [13]: # Split dataset into features (X) and target (y)  
X = iris_df.drop(['species', 'species_encoded'], axis=1)  
y = iris_df['species_encoded']
```

```
In [14]: # Split into training and testing sets  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

3. Feature Engineering

While the Iris dataset is relatively clean, you might want to scale the features for some models or engineer new features

```
In [15]: from sklearn.preprocessing import StandardScaler  
  
# Standardize features  
scaler = StandardScaler()  
X_train_scaled = scaler.fit_transform(X_train)  
X_test_scaled = scaler.transform(X_test)
```

4. Machine Learning Modeling Techniques

```
In [16]: from sklearn.svm import SVC  
from sklearn.metrics import classification_report, confusion_matrix
```

```
In [17]: # Initialize SVM classifier  
svm_clf = SVC(kernel='rbf', random_state=42)
```

```
In [18]: # Train the model  
svm_clf.fit(X_train_scaled, y_train)
```

```
Out[18]: SVC  
SVC(random_state=42)
```

```
In [19]: # Predictions  
y_pred = svm_clf.predict(X_test_scaled)
```

```
In [20]: # Evaluate the model  
print("Classification Report:")  
print(classification_report(y_test, y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	10
1	1.00	1.00	1.00	9
2	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

```
In [21]: # Confusion matrix
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, cmap='Blues', fmt='g', xticklabels=le.classes_, yticklabels=le.classes_,
            plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()
```

