

# 1. Data Analysis and Preprocessing

```
In [18]: # Import necessary Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
```

```
In [19]: # Load the dataset
data = pd.read_csv('advertising.csv') # replace with your actual data file
```

```
In [20]: # Display the first few rows of the dataset
print(data.head())
```

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9

```
In [21]: # Display summary statistics
print(data.describe())
```


	TV	Radio	Newspaper	Sales
count	200.000000	200.000000	200.000000	200.000000
mean	147.042500	23.264000	30.554000	15.130500
std	85.854236	14.846809	21.778621	5.283892
min	0.700000	0.000000	0.300000	1.600000
25%	74.375000	9.975000	12.750000	11.000000
50%	149.750000	22.900000	25.750000	16.000000
75%	218.825000	36.525000	45.100000	19.050000
max	296.400000	49.600000	114.000000	27.000000

```
In [22]: # Check for missing values
print(data.isnull().sum())
```


```
TV          0
Radio       0
Newspaper   0
Sales       0
dtype: int64
```

```
In [23]: # Visualize the relationships between features and target
sns.pairplot(data, x_vars=['TV', 'Radio', 'Newspaper'], y_vars='Sales', height=5,
plt.show())
```

```
C:\Users\Nishita Bala\anaconda3\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The figure layout has changed to tight
self._figure.tight_layout(*args, **kwargs)
```



```
In [24]: # Correlation matrix
corr_matrix = data.corr()
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm')
plt.show()
```



## 2. Feature Engineering and Data Splitting

```
In [25]: # Define features (X) and target (y)
X = data[['TV', 'Radio', 'Newspaper']]
y = data['Sales']
```

```
In [26]: # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [27]: # Standardize the features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

## 3. Modeling and Evaluation

### Linear Regression Model

```
In [28]: # Train the Linear Regression model
lr_model = LinearRegression()
lr_model.fit(X_train_scaled, y_train)
```

```
Out[28]: ▼ LinearRegression
LinearRegression()
```

```
In [29]: # Make predictions
y_pred_train = lr_model.predict(X_train_scaled)
y_pred_test = lr_model.predict(X_test_scaled)
```

```
In [30]: # Evaluate the model
train_rmse = np.sqrt(mean_squared_error(y_train, y_pred_train))
test_rmse = np.sqrt(mean_squared_error(y_test, y_pred_test))
train_r2 = r2_score(y_train, y_pred_train)
test_r2 = r2_score(y_test, y_pred_test)

print(f"Linear Regression Train RMSE: {train_rmse}")
print(f"Linear Regression Test RMSE: {test_rmse}")
print(f"Linear Regression Train R^2: {train_r2}")
print(f"Linear Regression Test R^2: {test_r2}")
```

```
Linear Regression Train RMSE: 1.6358920055378559
Linear Regression Test RMSE: 1.7052146229349232
Linear Regression Train R^2: 0.9001416005862131
Linear Regression Test R^2: 0.9059011844150826
```

## 4. Visualization

```
In [33]: # Visualize predictions vs actuals
plt.figure(figsize=(12, 6))

plt.subplot(1, 2, 1)
plt.scatter(y_test, y_pred_test, alpha=0.7)
plt.xlabel('Actual Sales')
plt.ylabel('Predicted Sales')

plt.tight_layout()
plt.show()
```

